

A decorative graphic at the top of the slide featuring a green sphere on the left and three overlapping semi-spheres (blue, red, and yellow) on the right.

ThreadSanitizer

Data race detection in practice

Konstantin Serebryany <kcc@google.com>

Timur Iskhodzhanov <timur.iskhodzhanov@phystech.edu>

Dec 12 2009



Data races are scary

A data race occurs when two or more threads concurrently access a shared memory location and at least one of the accesses is a write.

```
std::map<int,int> my_map;
```

```
void Thread1() {  
    my_map[123] = 1;  
}
```

```
void Thread2() {  
    my_map[345] = 2;  
}
```

Our goal: find races in Google code



Dynamic race detectors

- Intercept program events at run-time
 - Memory access: READ, WRITE
 - Synchronization: LOCK, UNLOCK, SIGNAL, WAIT
- Maintain global state
 - Locks, other synchronization events, threads
 - Memory allocations
- Maintain shadow state for each memory location (byte)
 - Remember previous accesses
 - Report race in appropriate state
- Two major approaches:
 - LockSet
 - Happens-before

LockSet


```
void Thread1() {  
    mu1.Lock();  
    mu2.Lock();  
    *X = 1;  
    mu2.Unlock();  
    mu1.Unlock(); ...  
}  
  
void Thread2() {  
    mu1.Lock();  
    mu3.Lock();  
    *X = 2;  
    mu3.Unlock();  
    mu1.Unlock(); ...  
}
```

- LockSet: a set of locks held during a memory access
 - Thread1: {mu1, mu2}
 - Thread2: {mu1, mu3}
- Common LockSet: intersection of LockSets
 - {mu1}

LockSet: false positives


```
void Thread1() {  
    x->Update(); // LS={}  
    mu.Lock();  
    queue.push(x);  
    mu.Unlock();  
}
```

```
void Thread2() {  
  
    Obj *y = NULL;  
    mu.Lock();  
    y = queue.pop_or_null();  
    mu.Unlock();  
    if (y) {  
        y->UseMe(); // LS={}  
    }  
}
```



Happens-before

partial order on all events



Segment: a sequence of READ/WRITE events of one thread
Signal(obj) \Rightarrow Wait(obj) is a happens-before arc

Seg1 $<$ Seg4 -- segments belong to the same thread.

Seg1 $<$ Seg5 -- due to Signal/Wait pair with a matching object.

Seg1 $<$ Seg7 -- happens-before is transitive.

Seg3 \nless Seg6 -- no ordering constraint.

Pure happens-before: misses races

```
void Thread1() {  
    x = 1;  
    mu.Lock();  
    do_something1();  
    mu.Unlock();  
}
```

```
void Thread2() {  
    do_something2();  
    // do_something2()  
    // may take  
    // lots of time  
    mu.Lock();  
    do_something3();  
    mu.Unlock();  
    x = 2;  
}
```

Happens-before vs LockSet

- Pure LockSet
 - Many false warnings (impractical in most cases)
 - Does not miss races, fast
- Pure happens-before detectors:
 - Unlock \Rightarrow Lock is a happens-before arc
 - No false positives
 - unless you use lock-free synchronization
 - Less predictable (bad for continuous build)
 - Misses many races (30% - 50%)
- Hybrid (LockSet for locks, happens-before for the rest):
 - No happens-before arc on Unlock/Lock
 - Has false positives (easy to annotate)
 - More predictable
 - Finds more races

Dynamic annotations

```
void Thread1() {  
    x->Update(); // LS={}  
    mu.Lock();  
    queue.push(x);  
    ANNOTATE_HAPPENS_BEFORE(x);  
    mu.Unlock();  
}
```

- Annotate lock-free synchronization
- Annotate false positives of hybrid mode
- Annotate benign races

```
void Thread2() {  
    Obj *y = NULL;  
    mu.Lock();  
    y = queue.pop_or_null();  
    ANNOTATE_HAPPENS_AFTER(y);  
    mu.Unlock();  
    if (y) {  
        y->UseMe(); // LS={}  
    }  
}
```

ThreadSanitizer: Algorithm

- Segment: a sequence of READ/WRITE events of one thread
 - All events in a segment have the same LockSet
- Segment Set: a set of segments none of which happen-before any other
- Shadow state:
 - Writer Segment Set: all recent writes
 - Reader Segment Set: all recent reads, unordered with or happened-after writes
- State machine: on each READ/WRITE event
 - update the Segment Sets
 - check accesses for race

ThreadSanitizer: Algorithm

```
HANDLE-READ-OR-WRITE-EVENT( $IsWrite, Tid, ID$ )  
1  ( $SS_{Wr}, SS_{Rd}$ )  $\leftarrow$  GET-PER-ID-STATE( $ID$ )  
2   $Seg \leftarrow$  GET-CURRENT-SEGMENT( $Tid$ )  
3  if  $IsWrite$   
4    then  $\triangleright$  WRITE event: update  $SS_{Wr}$  and  $SS_{Rd}$   
5           $SS_{Rd} \leftarrow \{s : s \in SS_{Rd} \wedge s \not\in Seg\}$   
6           $SS_{Wr} \leftarrow \{s : s \in SS_{Wr} \wedge s \not\in Seg\} \cup \{Seg\}$   
7    else  $\triangleright$  READ event: update  $SS_{Rd}$   
8           $SS_{Rd} \leftarrow \{s : s \in SS_{Rd} \wedge s \not\in Seg\} \cup \{Seg\}$   
9  SET-PER-ID-STATE( $ID, SS_{Wr}, SS_{Rd}$ )  
10 if IS-RACE( $SS_{Wr}, SS_{Rd}$ )  
11   then REPORT-RACE( $IsWrite, Tid, Seg, ID$ )
```

Example

// Thread1

X = ...;

Sem1.Post();

Sem2.Wait();

L1.Lock();

X = ...;

L1.Unlock();

// Thread2

Sem1.Wait();

... = X;

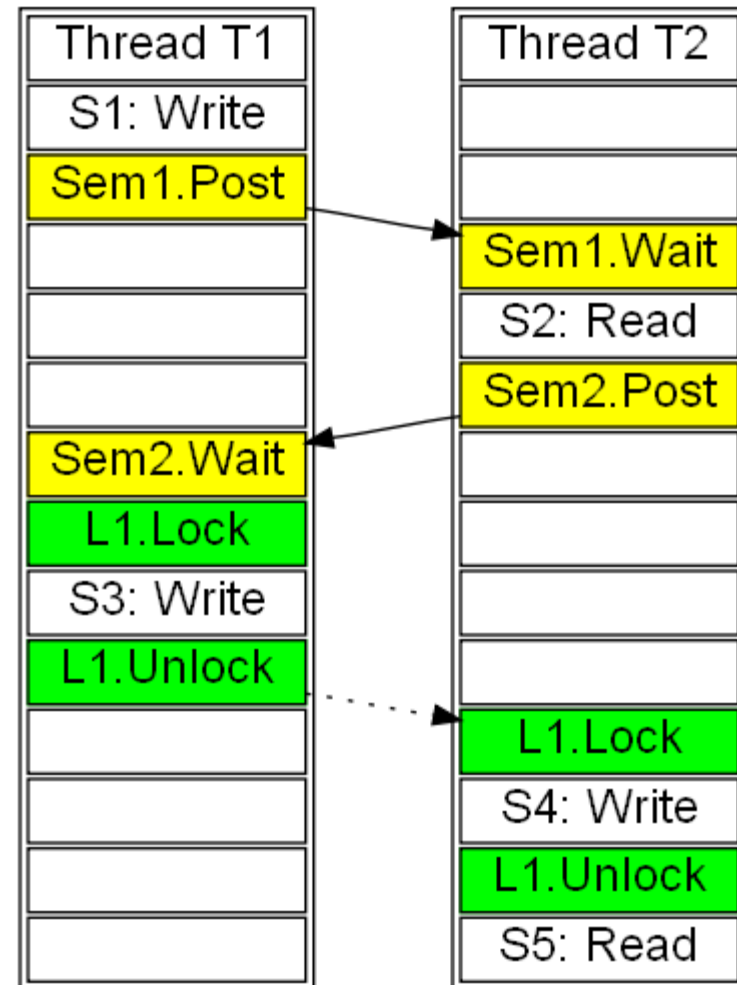
Sem2.Post();

L1.Lock();

X = ...;

L1.Unlock();

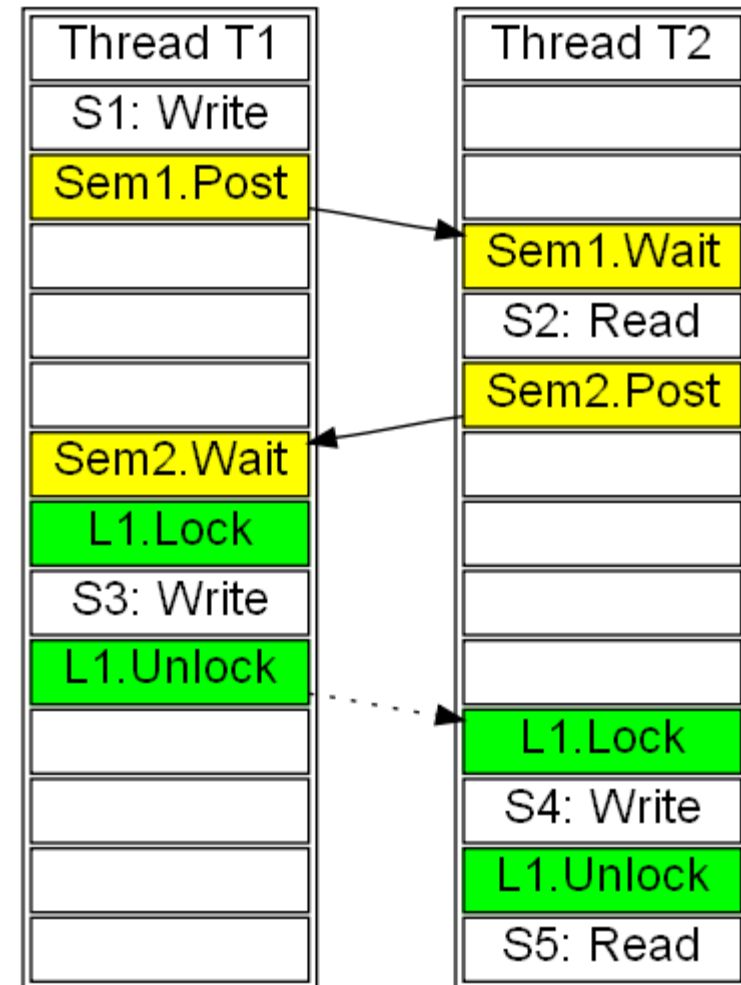
... = X;



Example: shadow state


Hybrid	
Writer SS	Reader SS
S1	-
S1	S2
S3/L1	-
S3/L1, S4/L1	-
S3/L1, S4/L1	S5 {Race!}

Pure Happens-before	
Writer SS	Reader SS
S1	-
S1	S2
S3/L1	-
S4/L1	-
S4/L1	S5




Report example

WARNING: Possible data race during write of size 4 at 0x633AA0: {{{


T2 (test-thread-2) (locks held: {L122}): 
#0 test301::Thread2() racecheck_unittest.cc:5956
#1 MyThread::ThreadBody(MyThread*) thread_wrappers_pthread.h:320
#2 ThreadSanitizerStartThread ts_valgrind_intercepts.c:387


Concurrent write(s) happened at (OR AFTER) these points:

T1 (test-thread-1) (locks held: {L121}): 
#0 test301::Thread1() racecheck_unittest.cc:5951
#1 MyThread::ThreadBody(MyThread*) thread_wrappers_pthread.h:320
#2 ThreadSanitizerStartThread ts_valgrind_intercepts.c:387

Address 0x633AA0 is 0 bytes inside data symbol "_ZN7test3013varE"

Locks involved in this report (reporting last lock sites): {L121, L122}

L121 (0x633A10) 
#0 pthread_mutex_lock ts_valgrind_intercepts.c:602
#1 Mutex::Lock() thread_wrappers_pthread.h:162
#2 MutexLock::MutexLock(Mutex*) thread_wrappers_pthread.h:225
#3 test301::Thread1() racecheck_unittest.cc:5950
#4 MyThread::ThreadBody(MyThread*) thread_wrappers_pthread.h:320
#5 ThreadSanitizerStartThread ts_valgrind_intercepts.c:387

L122 (0x633A70) 
#0 pthread_mutex_lock ts_valgrind_intercepts.c:602
#1 Mutex::Lock() thread_wrappers_pthread.h:162
#2 MutexLock::MutexLock(Mutex*) thread_wrappers_pthread.h:225
#3 test301::Thread2() racecheck_unittest.cc:5955
#4 MyThread::ThreadBody(MyThread*) thread_wrappers_pthread.h:320
#5 ThreadSanitizerStartThread ts_valgrind_intercepts.c:387

Conclusions

- Valgrind-based dynamic detector of data races for C/C++
- Works on Linux x86/x86_64 and Mac
- Has several modes of operation:
 - most conservative (pure happens-before):
 - few false reports, misses some races
 - most aggressive (hybrid):
 - more false positives, more real races
- Supports "Dynamic Annotations", a race detector API:
 - describe custom (e.g. lock-less) synchronization
 - hide benign races
 - zero noise level even in the most aggressive mode
- Opensource



Conclusions (cont)

- Overhead is comparable to Valgrind/Memcheck
 - Slowdown: 5x-50x
 - Memory overhead: 3x-6x
- Detailed output
 - Contains all locks involved in a race
- Runs regularly on thousands Google tests
 - Including all Chromium tests
- Found thousands of races
 - Several 'critical' (aka 'top crashers')
 - Dozens of potentially harmful
 - Tons of benign or test-only



A decorative header featuring four overlapping spheres: a green one on the left, and blue, red, and yellow ones on the right.

Q&A

<http://www.google.com/search?q=ThreadSanitizer&btnI>

